

Project 3: Memory Encrypting Disk Driver

Andrew Chase

OS II

Spring 2015

Abstract: Describes work done to build and run a custom Memory Encrypting Disk Driver.

CONTENTS

I	Work Log	3
II	Questions	3
II-A	The design you plan to use to implement the algorithms.	3
II-B	What do you think the main point of this assignment is?	3
II-C	How did you personally approach the problem? Design decisions, algorithm, etc. .	3
II-D	How did you ensure your solution was correct? Testing details, for instance. . . .	3
II-E	What did you learn?	4
III	Concurrency Exercise: Git Log	4
IV	Code Listing	4
IV-A	memdiskblk.c	4

I. WORK LOG

Date	Work Done
May 17	Write module, compile and test on os-class. Write helper scripts.
May 18	Fix bugs and test on os-class.

II. QUESTIONS

A. *The design you plan to use to implement the algorithms.*

The original plan was to use the cryptographic api exposed by Linux. The assignment gave an example device driver that contained the skeleton code needed to implement the memory disk device driver. The plan was to hook up the cryptographic primitives to that device driver such that data would be encrypted/decrypted in memory as requests were sent/received.

B. *What do you think the main point of this assignment is?*

I think the main point of the assignment is to use expose students to the block driver, crypto and other apis in the linux kernel and how they interact with modules. This also gives students a chance to learn about how to write custom drivers for the linux kernel. Another learning objective is exposure to the kernel memory api, which was discussed in class.

C. *How did you personally approach the problem? Design decisions, algorithm, etc.*

As before first I went out and sought other people's solutions to the problems. I read through them carefully and tried to figure out how they worked and what apis were used. I learned a lot from reading examples. Another thing I did was look deeply into the Linux kernel source code to try and figure out how the cryptographic primitives were supposed to be used.

I used the standard synchronous block api, with the one-block-at-a-time functions. The crypto function I used was AES.

D. *How did you ensure your solution was correct? Testing details, for instance.*

I inserted kprint calls and ran the kernel in the emulator. The kprint calls recorded the sector number.

Example:

```
Reading from memdiskblk --
```

```
Encrypted data: 000
```

```
Decrypted data: 001c
```

```
Reading from memdiskblk --
```

```
Encrypted data: 20ffffffa8ffffffc5
```

The testing helped because the first solution I had implemented wasn't correct and it was pretty obvious to see from the kernel messages. At first, my module didn't load, and using `dmesg` — `grep` I was able to find the error message to fix the problem.

E. What did you learn?

I learned a bit about how linux driver system works, how to compile and load modules into the linux kernel with parameters. I also learned that the linux kernel exposes a robust crypto api that can be easily integrated to add cryptographic primitives to drivers in the linux kernel. I also learned how to debug buggy kernel modules using `dmesg` and `gdb`.

III. CONCURRENCY EXERCISE: GIT LOG

acronym meaning

V	version
tag	git tag
MF	Number of modified files.
AL	Number of added lines.
DL	Number of deleted lines.

V	tag	date	commit message	MF	AL	DL
1		2016-05-18	init	0	191	0

IV. CODE LISTING

A. *memdiskblk.c*

```

1  /*
2   * Andy Chase
3   * CS444 -- Memory Encrypting Disk Driver
4   *
5   * Based on:
6   *
7   * http://blog.superpat.com/2010/05/04/a-simple-block-driver-for-linux-kernel-2-6-31/
8   * https://github.com/davidmerrick/Classes/blob/master/CS411/project3/osurd.c
9   * https://github.com/rleyherrington/linux_kernel_411/blob/master/device_driver/device_
10  * http://www.oreilly.com/openbook/linuxdrive3/book/ch16.pdf
11  * http://lxr.free-electrons.com/source/net/bluetooth/smp.c?v=3.14#L54

```

```

12  */
13
14  #include <linux/module.h>
15  #include <linux/moduleparam.h>
16  #include <linux/init.h>
17
18  #include <linux/kernel.h>          /* printk() */
19  #include <linux/fs.h>              /* everything... */
20  #include <linux/errno.h>          /* error codes */
21  #include <linux/types.h>          /* size_t */
22  #include <linux/vmalloc.h>
23  #include <linux/genhd.h>
24  #include <linux/blkdev.h>
25  #include <linux/hdreg.h>
26  #include <linux/crypto.h>
27  #include <linux/scatterlist.h>
28
29  /* Module Info */
30  MODULE_LICENSE("GPL");
31  MODULE_AUTHOR("Andrew Chase");
32  MODULE_DESCRIPTION("Homework 3: Encrypted Block Device");
33  MODULE_ALIAS("membdiskblk");
34
35  /* Parameters */
36  static int major_num = 0;
37  static int logical_block_size = 512;
38  static int nsectors = 1024;        /* How big the drive is */
39  static char *key = "some_key_yo";
40  module_param(key, charp, 0000);
41
42  #define KERNEL_SECTOR_SIZE 512
43
44  static struct request_queue *Queue;
45  struct crypto_cipher *tfm;

```

```

46
47 static struct memdiskblk_device {
48     unsigned long size;
49     spinlock_t lock;
50     u8 *data;
51     struct gendisk *gd;
52 } Device;
53
54 static void
55 memdiskblk_transfer(struct memdiskblk_device *dev, unsigned long sector,
56                    unsigned long nsect, char *buffer, int write)
57 {
58     unsigned long offset = sector * KERNEL_SECTOR_SIZE;
59     unsigned long nbytes = nsect * KERNEL_SECTOR_SIZE;
60     int i;
61     if ((offset + nbytes) > dev->size) {
62         printk(KERN_NOTICE "sbd: Beyond-end write (%ld %ld)\n", offset,
63                nbytes);
64         return;
65     }
66     crypto_cipher_setkey(tfm, key, strlen(key));
67
68     if (write) {
69         printk("Writing to memdiskblk -- \n");
70         if (nbytes > 3)
71             printk("Raw data: %x%x%x\n", (buffer)[0], (buffer)[1], (buffer)
72                for (i = 0; i < nbytes; i += crypto_cipher_blocksize(tfm)) {
73                 memset(dev->data + offset + i, 0,
74                        crypto_cipher_blocksize(tfm));
75                 crypto_cipher_encrypt_one(tfm, dev->data + offset + i,
76                        buffer + i);
77             }
78
79     if (nbytes > 3)

```

```

80         printk("Encrypted data: %x%x%x\n", (dev->data + offset)[0], (dev->data + offset)[1], (dev->data + offset)[2]);
81     } else {
82         printk("Reading from memdiskblk --\n");
83         if (nbytes > 3)
84             printk("Encrypted data: %x%x%x\n", (buffer)[0], (buffer)[1], (buffer)[2]);
85         for (i = 0; i < nbytes; i += crypto_cipher_blocksize(tfm)) {
86             crypto_cipher_decrypt_one(tfm, buffer + i,
87                                     dev->data + offset + i);
88         }
89         if (nbytes > 3)
90             printk("Decrypted data: %x%x%x\n", (dev->data + offset)[0], (dev->data + offset)[1], (dev->data + offset)[2]);
91     }
92 }
93
94 static void memdiskblk_request(struct request_queue *q)
95 {
96     struct request *req;
97
98     req = blk_fetch_request(q);
99     while (req != NULL) {
100         if (req == NULL || (req->cmd_type != REQ_TYPE_FS)) {
101             printk(KERN_NOTICE "Skip non-CMD request\n");
102             __blk_end_request_all(req, -EIO);
103             continue;
104         }
105         memdiskblk_transfer(&Device, blk_rq_pos(req),
106                           blk_rq_cur_sectors(req), req->buffer,
107                           rq_data_dir(req));
108         if (!__blk_end_request_cur(req, 0)) {
109             req = blk_fetch_request(q);
110         }
111     }
112 }
113

```

```

114 int memdiskblk_getgeo(struct block_device *block_device,
115                      struct hd_geometry *geo)
116 {
117     long size;
118
119     size = Device.size * (logical_block_size / KERNEL_SECTOR_SIZE);
120     geo->cylinders = (size & ~0x3f) >> 6;
121     geo->heads = 4;
122     geo->sectors = 16;
123     geo->start = 0;
124     return 0;
125 }
126
127 static struct block_device_operations memdiskblk_ops = {
128     .owner = THIS_MODULE,
129     .getgeo = memdiskblk_getgeo
130 };
131
132 static int __init memdiskblk_init(void)
133 {
134     tfm = crypto_alloc_cipher("aes", 0, 0);
135     /* Error checking for crypto */
136     if (IS_ERR(tfm)) {
137         printk(KERN_ERR "memdiskblk -- cipher allocation failed");
138         return PTR_ERR(tfm);
139     }
140
141     Device.size = nsectors * logical_block_size;
142     spin_lock_init(&Device.lock);
143     Device.data = vmalloc(Device.size);
144     if (Device.data == NULL)
145         return -ENOMEM;
146
147     Queue = blk_init_queue(memdiskblk_request, &Device.lock);

```



```

148     if (Queue == NULL)
149         goto out;
150     blk_queue_logical_block_size(Queue, logical_block_size);
151
152     major_num = register_blkdev(major_num, "memdiskblk");
153     if (major_num < 0) {
154         printk(KERN_WARNING "memdiskblk: unable to get major number\n");
155         goto out;
156     }
157
158     Device.gd = alloc_disk(16);
159     if (!Device.gd)
160         goto out_unregister;
161     Device.gd->major = major_num;
162     Device.gd->first_minor = 0;
163     Device.gd->fops = &memdiskblk_ops;
164     Device.gd->private_data = &Device;
165     strcpy(Device.gd->disk_name, "memdiskblk0");
166     set_capacity(Device.gd, nsectors);
167     Device.gd->queue = Queue;
168     add_disk(Device.gd);
169
170     return 0;
171
172 out_unregister:
173     unregister_blkdev(major_num, "memdiskblk");
174 out:
175     crypto_free_cipher(tfm);
176     vfree(Device.data);
177     return -ENOMEM;
178 }
179
180 static void __exit memdiskblk_exit(void)
181 {

```

```
182         del_gendisk (Device.gd);
183         put_disk (Device.gd);
184         unregister_blkdev (major_num, "memdiskblk");
185         blk_cleanup_queue (Queue);
186         vfree (Device.data);
187     }
188
189     module_init (memdiskblk_init);
190     module_exit (memdiskblk_exit);
```