

Project 4: The SLOB SLAB

Andrew Chase

OS II

Spring 2015

Abstract: Describes work done to build and run a custom Memory Allocator.

CONTENTS

I	Work Log	3
II	Questions	3
II-A	What do you think the main point of this assignment is?	3
II-B	How did you personally approach the problem? Design decisions, algorithm, etc. . .	3
II-C	How did you ensure your solution was correct? Testing details, for instance.	3
II-D	What did you learn?	3
III	Git Log	3
IV	Code Listing	4
IV-A	slob.c	4
IV-B	frag.py	23

I. WORK LOG

Date Work Done

June 3 Did research, modified slob.c, wrote frag.py

II. QUESTIONS

A. What do you think the main point of this assignment is?

I think the main point of the assignment is to learn about memory management in operating systems and how the Linux memory management system works.

B. How did you personally approach the problem? Design decisions, algorithm, etc.

As usual first I looked up other approaches to the problem. Next I designed a simple algorithm. I focused on trying to make as few changes to the slob.c module as possible. My algorithm is a little wasteful as it goes through the whole list of blocks twice, once to find the best fit, and once to re-find the best-fit block to select it for use.

For the memory fragmentation I found a reference for a linux system tool that allows me to grab the list of blocks easily so that I can perform operations on them.

C. How did you ensure your solution was correct? Testing details, for instance.

I ensured my solution was correct by reading the `/proc/buddyinfo` file which reveals the free memory blocks and their sizes. I also inserted `printk` messages to example the values of the variables at one point when my solution wasn't working correctly.

D. What did you learn?

I learned a bit about how the Linux block allocator works and how the different algorithms can effect performance. I also learned a bit about how to debug memory errors in QEMU. I learned how to perform research and wire up solutions to problems in Kernel development.

III. GIT LOG

acronym	meaning
V	version
tag	git tag
MF	Number of modified files.
AL	Number of added lines.
DL	Number of deleted lines.

V	tag	date	commit message	MF	AL	DL
1		2014-12-09	Merge tag 'v3.14.24'	45949	18281318	0
2		2016-06-03	Add plan	1	14	0
3		2016-06-03	Implement best-first	1	27	30
4		2016-06-03	Fix int declarations to c90	1	5	4
5		2016-06-03	Fix additional warnings	1	7	5
6		2016-06-03	Fix bug where slob would never allocate	1	3	3
7		2016-06-03	Fix reversed for loop args	1	1	1
8		2016-06-03	Fix base case	1	1	1
9		2016-06-03	Fix no memory available	1	4	2
10		2016-06-03	Fix memory freeze	1	5	2
11		2016-06-03	Add fragmentation percentage script	1	32	0

IV. CODE LISTING

A. *slob.c*

```

1  /*
2   * SLOB Allocator: Simple List Of Blocks
3   *
4   * Matt Mackall <mpm@selenic.com> 12/30/03
5   *
6   * NUMA support by Paul Mundt, 2007.
7   *
8   * How SLOB works:
9   *
10  * The core of SLOB is a traditional K&R style heap allocator, with
11  * support for returning aligned objects. The granularity of this
12  * allocator is as little as 2 bytes, however typically most architectures
13  * will require 4 bytes on 32-bit and 8 bytes on 64-bit.
14  *
15  * The slob heap is a set of linked list of pages from alloc_pages(),
16  * and within each page, there is a singly-linked list of free blocks
17  * (slob_t). The heap is grown on demand. To reduce fragmentation,
18  * heap pages are segregated into three lists, with objects less than

```

```
19  * 256 bytes, objects less than 1024 bytes, and all other objects.
20  *
21  * Allocation from heap involves first searching for a page with
22  * sufficient free blocks (using a next-fit-like approach) followed by
23  * a first-fit scan of the page. Deallocation inserts objects back
24  * into the free list in address order, so this is effectively an
25  * address-ordered first fit.
26  *
27  * Above this is an implementation of kmalloc/kfree. Blocks returned
28  * from kmalloc are prepended with a 4-byte header with the kmalloc size.
29  * If kmalloc is asked for objects of PAGE_SIZE or larger, it calls
30  * alloc_pages() directly, allocating compound pages so the page order
31  * does not have to be separately tracked.
32  * These objects are detected in kfree() because PageSlab()
33  * is false for them.
34  *
35  * SLAB is emulated on top of SLOB by simply calling constructors and
36  * destructors for every SLAB allocation. Objects are returned with the
37  * 4-byte alignment unless the SLAB_HWCACHE_ALIGN flag is set, in which
38  * case the low-level allocator will fragment blocks to create the proper
39  * alignment. Again, objects of page-size or greater are allocated by
40  * calling alloc_pages(). As SLAB objects know their size, no separate
41  * size bookkeeping is necessary and there is essentially no allocation
42  * space overhead, and compound pages aren't needed for multi-page
43  * allocations.
44  *
45  * NUMA support in SLOB is fairly simplistic, pushing most of the real
46  * logic down to the page allocator, and simply doing the node accounting
47  * on the upper levels. In the event that a node id is explicitly
48  * provided, alloc_pages_exact_node() with the specified node id is used
49  * instead. The common case (or when the node id isn't explicitly provided)
50  * will default to the current node, as per numa_node_id().
51  *
52  * Node aware pages are still inserted in to the global freelist, and
```

```

53  * these are scanned for by matching against the node id encoded in the
54  * page flags. As a result, block allocations that can be satisfied from
55  * the freelist will only be done so on pages residing on the same node,
56  * in order to prevent random node placement.
57  */
58
59  #include <linux/kernel.h>
60  #include <linux/slab.h>
61
62  #include <linux/mm.h>
63  #include <linux/swap.h> /* struct reclaim_state */
64  #include <linux/cache.h>
65  #include <linux/init.h>
66  #include <linux/export.h>
67  #include <linux/rcupdate.h>
68  #include <linux/list.h>
69  #include <linux/kmemleak.h>
70
71  #include <trace/events/kmem.h>
72
73  #include <linux/atomic.h>
74
75  #include "slab.h"
76  /*
77   * slob_block has a field 'units', which indicates size of block if +ve,
78   * or offset of next block if -ve (in SLOB_UNITS).
79   *
80   * Free blocks of size 1 unit simply contain the offset of the next block.
81   * Those with larger size contain their size in the first SLOB_UNIT of
82   * memory, and the offset of the next free block in the second SLOB_UNIT.
83   */
84  #if PAGE_SIZE <= (32767 * 2)
85  typedef s16 slobidx_t;
86  #else

```

```

87 typedef s32 slobidx_t;
88 #endif
89
90 struct slob_block {
91     slobidx_t units;
92 };
93 typedef struct slob_block slob_t;
94
95 /*
96  * All partially free slob pages go on these lists.
97  */
98 #define SLOB_BREAK1 256
99 #define SLOB_BREAK2 1024
100 static LIST_HEAD(free_slob_small);
101 static LIST_HEAD(free_slob_medium);
102 static LIST_HEAD(free_slob_large);
103
104 /*
105  * slob_page_free: true for pages on free_slob_pages list.
106  */
107 static inline int slob_page_free(struct page *sp)
108 {
109     return PageSlobFree(sp);
110 }
111
112 static void set_slob_page_free(struct page *sp, struct list_head *list)
113 {
114     list_add(&sp->list, list);
115     __SetPageSlobFree(sp);
116 }
117
118 static inline void clear_slob_page_free(struct page *sp)
119 {
120     list_del(&sp->list);

```

```

121     __ClearPageSlobFree(sp);
122 }
123
124 #define SLOB_UNIT sizeof(slob_t)
125 #define SLOB_UNITS(size) DIV_ROUND_UP(size, SLOB_UNIT)
126
127 /*
128  * struct slob_rcu is inserted at the tail of allocated slob blocks, which
129  * were created with a SLAB_DESTROY_BY_RCU slab. slob_rcu is used to free
130  * the block using call_rcu.
131  */
132 struct slob_rcu {
133     struct rcu_head head;
134     int size;
135 };
136
137 /*
138  * slob_lock protects all slob allocator structures.
139  */
140 static DEFINE_SPINLOCK(slob_lock);
141
142 /*
143  * Encode the given size and next info into a free slob block s.
144  */
145 static void set_slob(slob_t *s, slobidx_t size, slob_t *next)
146 {
147     slob_t *base = (slob_t *) ((unsigned long)s & PAGE_MASK);
148     slobidx_t offset = next - base;
149
150     if (size > 1) {
151         s[0].units = size;
152         s[1].units = offset;
153     } else
154         s[0].units = -offset;

```



```

155 }
156
157 /*
158  * Return the size of a slob block.
159  */
160 static slobidx_t slob_units(slob_t *s)
161 {
162     if (s->units > 0)
163         return s->units;
164     return 1;
165 }
166
167 /*
168  * Return the next free slob block pointer after this one.
169  */
170 static slob_t *slob_next(slob_t *s)
171 {
172     slob_t *base = (slob_t *) ((unsigned long) s & PAGE_MASK);
173     slobidx_t next;
174
175     if (s[0].units < 0)
176         next = -s[0].units;
177     else
178         next = s[1].units;
179     return base+next;
180 }
181
182 /*
183  * Returns true if s is the last free block in its page.
184  */
185 static int slob_last(slob_t *s)
186 {
187     return !((unsigned long) slob_next(s) & ~PAGE_MASK);
188 }

```

```

189
190 static void *slob_new_pages(gfp_t gfp, int order, int node)
191 {
192     void *page;
193
194     #ifdef CONFIG_NUMA
195         if (node != NUMA_NO_NODE)
196             page = alloc_pages_exact_node(node, gfp, order);
197         else
198             #endif
199             page = alloc_pages(gfp, order);
200
201     if (!page)
202         return NULL;
203
204     return page_address(page);
205 }
206
207 static void slob_free_pages(void *b, int order)
208 {
209     if (current->reclaim_state)
210         current->reclaim_state->reclaimed_slab += 1 << order;
211     free_pages((unsigned long)b, order);
212 }
213
214 /*
215  * Allocate a slob block within a given slob_page sp.
216  */
217 static void *slob_page_alloc(struct page *sp, size_t size, int align, int apply, int *b
218 {
219     slob_t *prev, *cur, *aligned = NULL;
220     int delta = 0, units = SLOB_UNITS(size);
221     int fits;
222

```

```

223     for (prev = NULL, cur = sp->freelist; ; prev = cur, cur = slob_next(cur)) {
224         slobidx_t avail = slob_units(cur);
225
226         if (align) {
227             aligned = (slob_t *)ALIGN((unsigned long)cur, align);
228             delta = aligned - cur;
229         }
230         fits = avail >= units + delta;
231         if (!apply && fits && (!*best_fit || (avail - units + delta < *best_fit)))
232             *best_fit = avail - units + delta;
233     }
234
235     if (apply && fits && (!*best_fit || avail - units + delta == *best_fit))
236         slob_t *next;
237
238         if (delta) { /* need to fragment head to align? */
239             next = slob_next(cur);
240             set_slob(aligned, avail - delta, next);
241             set_slob(cur, delta, aligned);
242             prev = cur;
243             cur = aligned;
244             avail = slob_units(cur);
245         }
246
247         next = slob_next(cur);
248         if (avail == units) { /* exact fit? unlink. */
249             if (prev)
250                 set_slob(prev, slob_units(prev), next);
251             else
252                 sp->freelist = next;
253         } else { /* fragment */
254             if (prev)
255                 set_slob(prev, slob_units(prev), cur + units);
256             else

```

```

257         sp->freelist = cur + units;
258         set_slob(cur + units, avail - units, next);
259     }
260
261     sp->units -= units;
262     if (!sp->units)
263         clear_slob_page_free(sp);
264     return cur;
265 }
266 if (slob_last(cur))
267     return NULL;
268 }
269 }
270
271 /*
272  * slob_alloc: entry point into the slob allocator.
273  */
274 static void *slob_alloc(size_t size, gfp_t gfp, int align, int node)
275 {
276     struct page *sp;
277     struct list_head *prev;
278     struct list_head *slob_list;
279     slob_t *b = NULL;
280     unsigned long flags;
281     int best_fit = 0;
282     int apply;
283
284     if (size < SLOB_BREAK1)
285         slob_list = &free_slob_small;
286     else if (size < SLOB_BREAK2)
287         slob_list = &free_slob_medium;
288     else
289         slob_list = &free_slob_large;
290

```

```

291 spin_lock_irqsave(&slob_lock, flags);
292 /* Iterate through each partially free page, try to find room */
293 for(apply = 0; apply <= 1; apply++) {
294     list_for_each_entry(sp, slob_list, list) {
295         #ifdef CONFIG_NUMA
296             /*
297              * If there's a node specification, search for a partial
298              * page with a matching node id in the freelist.
299              */
300             if (node != NUMA_NO_NODE && page_to_nid(sp) != node)
301                 continue;
302         #endif
303         /* Enough room on this page? */
304         if (sp->units < SLOB_UNITS(size))
305             continue;
306
307         /* Attempt to alloc */
308         prev = sp->list.prev;
309         b = slob_page_alloc(sp, size, align, apply, &best_fit);
310         if (!b)
311             continue;
312         break;
313     }
314 }
315 spin_unlock_irqrestore(&slob_lock, flags);
316
317 /* Not enough space: must allocate a new page */
318 if (!b) {
319     b = slob_new_pages(gfp & ~__GFP_ZERO, 0, node);
320     if (!b)
321         return NULL;
322     sp = virt_to_page(b);
323     __SetPageSlab(sp);
324

```

```

325         spin_lock_irqsave(&slob_lock, flags);
326         sp->units = SLOB_UNITS(PAGE_SIZE);
327         sp->freelist = b;
328         INIT_LIST_HEAD(&sp->list);
329         set_slob(b, SLOB_UNITS(PAGE_SIZE), b + SLOB_UNITS(PAGE_SIZE));
330         set_slob_page_free(sp, slob_list);
331         best_fit = 0;
332         apply = 1;
333         b = slob_page_alloc(sp, size, align, apply, &best_fit);
334         BUG_ON(!b);
335         spin_unlock_irqrestore(&slob_lock, flags);
336     }
337     if (unlikely((gfp & __GFP_ZERO) && b))
338         memset(b, 0, size);
339     return b;
340 }
341
342 /*
343  * slob_free: entry point into the slob allocator.
344  */
345 static void slob_free(void *block, int size)
346 {
347     struct page *sp;
348     slob_t *prev, *next, *b = (slob_t *)block;
349     slobidx_t units;
350     unsigned long flags;
351     struct list_head *slob_list;
352
353     if (unlikely(ZERO_OR_NULL_PTR(block)))
354         return;
355     BUG_ON(!size);
356
357     sp = virt_to_page(block);
358     units = SLOB_UNITS(size);

```

```

359
360     spin_lock_irqsave(&slob_lock, flags);
361
362     if (sp->units + units == SLOB_UNITS(PAGE_SIZE)) {
363         /* Go directly to page allocator. Do not pass slob allocator */
364         if (slob_page_free(sp))
365             clear_slob_page_free(sp);
366         spin_unlock_irqrestore(&slob_lock, flags);
367         __ClearPageSlab(sp);
368         page_mapcount_reset(sp);
369         slob_free_pages(b, 0);
370         return;
371     }
372
373     if (!slob_page_free(sp)) {
374         /* This slob page is about to become partially free. Easy! */
375         sp->units = units;
376         sp->freelist = b;
377         set_slob(b, units,
378                 (void *) ((unsigned long) (b +
379                                         SLOB_UNITS(PAGE_SIZE)) & PAGE_MASK));
380         if (size < SLOB_BREAK1)
381             slob_list = &free_slob_small;
382         else if (size < SLOB_BREAK2)
383             slob_list = &free_slob_medium;
384         else
385             slob_list = &free_slob_large;
386         set_slob_page_free(sp, slob_list);
387         goto out;
388     }
389
390     /*
391      * Otherwise the page is already partially free, so find reinsertion
392      * point.

```

```

393     */
394     sp->units += units;
395
396     if (b < (slob_t *)sp->freelist) {
397         if (b + units == sp->freelist) {
398             units += slob_units(sp->freelist);
399             sp->freelist = slob_next(sp->freelist);
400         }
401         set_slob(b, units, sp->freelist);
402         sp->freelist = b;
403     } else {
404         prev = sp->freelist;
405         next = slob_next(prev);
406         while (b > next) {
407             prev = next;
408             next = slob_next(prev);
409         }
410
411         if (!slob_last(prev) && b + units == next) {
412             units += slob_units(next);
413             set_slob(b, units, slob_next(next));
414         } else
415             set_slob(b, units, next);
416
417         if (prev + slob_units(prev) == b) {
418             units = slob_units(b) + slob_units(prev);
419             set_slob(prev, units, slob_next(b));
420         } else
421             set_slob(prev, slob_units(prev), b);
422     }
423 out:
424     spin_unlock_irqrestore(&slob_lock, flags);
425 }
426

```



```

427  /*
428   * End of slob allocator proper. Begin kmem_cache_alloc and kmalloc frontend.
429   */
430
431  static __always_inline void *
432  __do_kmalloc_node(size_t size, gfp_t gfp, int node, unsigned long caller)
433  {
434      unsigned int *m;
435      int align = max_t(size_t, ARCH_KMALLOC_MINALIGN, ARCH_SLAB_MINALIGN);
436      void *ret;
437
438      gfp &= gfp_allowed_mask;
439
440      lockdep_trace_alloc(gfp);
441
442      if (size < PAGE_SIZE - align) {
443          if (!size)
444              return ZERO_SIZE_PTR;
445
446          m = slob_alloc(size + align, gfp, align, node);
447
448          if (!m)
449              return NULL;
450
451          *m = size;
452          ret = (void *)m + align;
453
454          trace_kmalloc_node(caller, ret,
455                           size, size + align, gfp, node);
456      } else {
457          unsigned int order = get_order(size);
458
459          if (likely(order))
460              gfp |= __GFP_COMP;
461          ret = slob_new_pages(gfp, order, node);

```

```

461
462         trace_kmalloc_node(caller, ret,
463                             size, PAGE_SIZE << order, gfp, node);
464     }
465
466     kmemleak_alloc(ret, size, 1, gfp);
467     return ret;
468 }
469
470 void *__kmalloc(size_t size, gfp_t gfp)
471 {
472     return __do_kmalloc_node(size, gfp, NUMA_NO_NODE, _RET_IP_);
473 }
474 EXPORT_SYMBOL(__kmalloc);
475
476 #ifdef CONFIG_TRACING
477 void *__kmalloc_track_caller(size_t size, gfp_t gfp, unsigned long caller)
478 {
479     return __do_kmalloc_node(size, gfp, NUMA_NO_NODE, caller);
480 }
481
482 #ifdef CONFIG_NUMA
483 void *__kmalloc_node_track_caller(size_t size, gfp_t gfp,
484                                   int node, unsigned long caller)
485 {
486     return __do_kmalloc_node(size, gfp, node, caller);
487 }
488 #endif
489 #endif
490
491 void kfree(const void *block)
492 {
493     struct page *sp;
494

```

```

495     trace_kfree(_RET_IP_, block);
496
497     if (unlikely(ZERO_OR_NULL_PTR(block)))
498         return;
499     kmemleak_free(block);
500
501     sp = virt_to_page(block);
502     if (PageSlab(sp)) {
503         int align = max_t(size_t, ARCH_KMALLOC_MINALIGN, ARCH_SLAB_MINALIGN);
504         unsigned int *m = (unsigned int *) (block - align);
505         slob_free(m, *m + align);
506     } else
507         __free_pages(sp, compound_order(sp));
508 }
509 EXPORT_SYMBOL(kfree);
510
511 /* can't use ksize for kmem_cache_alloc memory, only kmalloc */
512 size_t ksize(const void *block)
513 {
514     struct page *sp;
515     int align;
516     unsigned int *m;
517
518     BUG_ON(!block);
519     if (unlikely(block == ZERO_SIZE_PTR))
520         return 0;
521
522     sp = virt_to_page(block);
523     if (unlikely(!PageSlab(sp)))
524         return PAGE_SIZE << compound_order(sp);
525
526     align = max_t(size_t, ARCH_KMALLOC_MINALIGN, ARCH_SLAB_MINALIGN);
527     m = (unsigned int *) (block - align);
528     return SLOB_UNITS(*m) * SLOB_UNIT;

```

```

529 }
530 EXPORT_SYMBOL(ksize);
531
532 int __kmem_cache_create(struct kmem_cache *c, unsigned long flags)
533 {
534     if (flags & SLAB_DESTROY_BY_RCU) {
535         /* leave room for rcu footer at the end of object */
536         c->size += sizeof(struct slob_rcu);
537     }
538     c->flags = flags;
539     return 0;
540 }
541
542 void *slob_alloc_node(struct kmem_cache *c, gfp_t flags, int node)
543 {
544     void *b;
545
546     flags &= gfp_allowed_mask;
547
548     lockdep_trace_alloc(flags);
549
550     if (c->size < PAGE_SIZE) {
551         b = slob_alloc(c->size, flags, c->align, node);
552         trace_kmem_cache_alloc_node(_RET_IP_, b, c->object_size,
553                                     SLOB_UNITS(c->size) * SLOB_UNIT,
554                                     flags, node);
555     } else {
556         b = slob_new_pages(flags, get_order(c->size), node);
557         trace_kmem_cache_alloc_node(_RET_IP_, b, c->object_size,
558                                     PAGE_SIZE << get_order(c->size),
559                                     flags, node);
560     }
561
562     if (b && c->ctor)

```

```

563         c->ctor(b);
564
565         kmemleak_alloc_recursive(b, c->size, 1, c->flags, flags);
566         return b;
567     }
568     EXPORT_SYMBOL(slob_alloc_node);
569
570     void *kmem_cache_alloc(struct kmem_cache *cachep, gfp_t flags)
571     {
572         return slob_alloc_node(cachep, flags, NUMA_NO_NODE);
573     }
574     EXPORT_SYMBOL(kmem_cache_alloc);
575
576     #ifdef CONFIG_NUMA
577     void *__kmalloc_node(size_t size, gfp_t gfp, int node)
578     {
579         return __do_kmalloc_node(size, gfp, node, _RET_IP_);
580     }
581     EXPORT_SYMBOL(__kmalloc_node);
582
583     void *kmem_cache_alloc_node(struct kmem_cache *cachep, gfp_t gfp, int node)
584     {
585         return slob_alloc_node(cachep, gfp, node);
586     }
587     EXPORT_SYMBOL(kmem_cache_alloc_node);
588     #endif
589
590     static void __kmem_cache_free(void *b, int size)
591     {
592         if (size < PAGE_SIZE)
593             slob_free(b, size);
594         else
595             slob_free_pages(b, get_order(size));
596     }

```

```

597
598 static void kmem_rcu_free(struct rcu_head *head)
599 {
600     struct slob_rcu *slob_rcu = (struct slob_rcu *)head;
601     void *b = (void *)slob_rcu - (slob_rcu->size - sizeof(struct slob_rcu));
602
603     __kmem_cache_free(b, slob_rcu->size);
604 }
605
606 void kmem_cache_free(struct kmem_cache *c, void *b)
607 {
608     kmemleak_free_recursive(b, c->flags);
609     if (unlikely(c->flags & SLAB_DESTROY_BY_RCU)) {
610         struct slob_rcu *slob_rcu;
611         slob_rcu = b + (c->size - sizeof(struct slob_rcu));
612         slob_rcu->size = c->size;
613         call_rcu(&slob_rcu->head, kmem_rcu_free);
614     } else {
615         __kmem_cache_free(b, c->size);
616     }
617
618     trace_kmem_cache_free(_RET_IP_, b);
619 }
620 EXPORT_SYMBOL(kmem_cache_free);
621
622 int __kmem_cache_shutdown(struct kmem_cache *c)
623 {
624     /* No way to check for remaining objects */
625     return 0;
626 }
627
628 int kmem_cache_shrink(struct kmem_cache *d)
629 {
630     return 0;

```

```

631 }
632 EXPORT_SYMBOL(kmem_cache_shrink);
633
634 struct kmem_cache kmem_cache_boot = {
635     .name = "kmem_cache",
636     .size = sizeof(struct kmem_cache),
637     .flags = SLAB_PANIC,
638     .align = ARCH_KMALLOC_MINALIGN,
639 };
640
641 void __init kmem_cache_init(void)
642 {
643     kmem_cache = &kmem_cache_boot;
644     slab_state = UP;
645 }
646
647 void __init kmem_cache_init_late(void)
648 {
649     slab_state = FULL;
650 }

```

B. frag.py

```

1  #!/usr/bin/python
2  import sys
3
4  mult_list = [
5      4,
6      8,
7      16,
8      32,
9      64,
10     128,
11     256,
12     512,

```

```

13     1024,
14     2048,
15     4096
16 ]
17
18 example_data = "Node 0, zone    Normal    75    51    57    16    2    0    0
19 example_data_2 = "Normal    272    144    80    50    33    13    2    1    2
20
21 def get_number_list(input_string):
22     return (int(i) for i in input_string.split("zone    Normal")[1].split())
23
24 def get_fragmentation(input_blocks):
25     return (sum(input_blocks) - input_blocks[-1]) / float(sum(input_blocks))
26
27 def size_blocks(number_list):
28     for block, size in zip(number_list, mult_list):
29         yield block*size
30
31 data = open("/proc/buddyinfo", "r").read()
32 print(get_fragmentation(list(size_blocks(get_number_list(data)))))

```